

```
# from sklearn.datasets import load_wine
# data = load_wine()
# X, y = data.data, data.target

# from sklearn.datasets import load_digits
# data = load_digits()
# X, y = data.data, data.target

from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split

data = load_digits()
X, y = data.data, data.target

#splits = [0.5, 0.4, 0.3, 0.2]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix

kernels = ['linear', 'poly', 'rbf', 'sigmoid']
for k in kernels:
    model = SVC(kernel=k)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='macro') # precision
    rec = recall_score(y_test, y_pred, average='macro') # recall
    f1 = f1_score(y_test, y_pred, average='macro') # f1-score
    cm = confusion_matrix(y_test, y_pred)

    print(f"Kernel: {k}")
    print(f"Accuracy : {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall : {rec:.4f}")
    print(f"F1-score : {f1:.4f}")
    print("Confusion Matrix:\n", cm)
    print("-"*40)
```

```
Kernel: linear
Accuracy : 0.9759
Precision: 0.9761
Recall : 0.9755
F1-score : 0.9755
Confusion Matrix:
[[53  0  0  0  1  0  0  0  0  0]
 [ 0 54  0  0  0  0  0  0  0  1]
 [ 0  0 53  0  0  0  0  0  0  0]
 [ 0  0  0 54  0  0  0  0  0  1]
 [ 0  0  0  0 54  0  0  0  0  0]
 [ 0  0  0  0  0 54  0  0  0  1]
 [ 0  0  0  0  0  0 54  0  0  0]
 [ 0  0  0  0  0  0  0 54  0  0]
 [ 0  4  0  0  0  1  1  1 45  0]
 [ 0  0  0  0  0  0  0  0  2 52]]
```

```
-----
Kernel: poly
Accuracy : 0.9889
Precision: 0.9890
Recall : 0.9888
F1-score : 0.9888
Confusion Matrix:
[[53  0  0  0  1  0  0  0  0  0]
 [ 0 55  0  0  0  0  0  0  0  0]
 [ 0  0 53  0  0  0  0  0  0  0]
 [ 0  0  0 55  0  0  0  0  0  0]
 [ 0  0  0  0 54  0  0  0  0  0]
 [ 0  0  0  0  0 54  0  0  0  1]
 [ 0  0  0  0  0  0 53  0  1  0]
```

```
[ 0 0 0 0 0 0 0 0 54 0 0]
[ 0 2 0 0 0 0 0 0 50 0]
[ 0 0 0 0 0 0 0 1 0 53]]
```

Kernel: rbf

Accuracy : 0.9889

Precision: 0.9890

Recall : 0.9888

F1-score : 0.9888

Confusion Matrix:

```
[[53 0 0 0 1 0 0 0 0 0]
 [ 0 55 0 0 0 0 0 0 0 0]
 [ 0 0 53 0 0 0 0 0 0 0]
 [ 0 0 0 55 0 0 0 0 0 0]
 [ 0 0 0 0 53 0 0 0 1 0]
 [ 0 0 0 0 0 54 0 0 0 1]
 [ 0 0 0 0 0 0 54 0 0 0]
 [ 0 0 0 0 0 0 0 54 0 0]
 [ 0 2 0 0 0 0 0 0 50 0]
 [ 0 0 0 0 0 0 0 1 0 53]]
```

Kernel: sigmoid

Accuracy : 0.8944

Precision: 0.9010

Recall : 0.8936

F1-score : 0.8937

Confusion Matrix:

```
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
```

```
data = load_digits()
X, y = data.data, data.target
```

```
#splits = [0.5, 0.4, 0.3, 0.2]
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
```

```
kernels = ['linear', 'poly', 'rbf', 'sigmoid']
```

```
for k in kernels:
```

```
    model = SVC(kernel=k)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
```

```
    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='macro') # precision
    rec = recall_score(y_test, y_pred, average='macro') # recall
    f1 = f1_score(y_test, y_pred, average='macro') # f1-score
    cm = confusion_matrix(y_test, y_pred)
```

```
    print(f"Kernel: {k}")
    print(f"Accuracy : {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall : {rec:.4f}")
    print(f"F1-score : {f1:.4f}")
    print("Confusion Matrix:\n", cm)
    print("-"*40)
```



Kernel: linear

Accuracy : 0.9778

Precision: 0.9779

Recall : 0.9775

F1-score : 0.9775

Confusion Matrix:

```
[[36 0 0 0 0 0 0 0 0 0]
 [ 0 34 0 0 0 0 0 0 1 1]
 [ 0 0 35 0 0 0 0 0 0 0]
 [ 0 0 0 36 0 0 0 0 0 1]
 [ 0 0 0 0 36 0 0 0 0 0]
 [ 0 0 0 0 0 37 0 0 0 0]
 [ 0 0 0 0 0 0 35 0 1 0]
 [ 0 0 0 0 0 0 0 36 0 0]
 [ 0 3 0 0 0 0 0 1 31 0]
 [ 0 0 0 0 0 0 0 0 0 36]]
```

Kernel: poly

Accuracy : 0.9861

Precision: 0.9864

```

Recall    : 0.9860
F1-score  : 0.9860
Confusion Matrix:
[[35  0  0  0  1  0  0  0  0  0]
 [ 0 36  0  0  0  0  0  0  0  0]
 [ 0  0 35  0  0  0  0  0  0  0]
 [ 0  0  0 37  0  0  0  0  0  0]
 [ 0  0  0  0 36  0  0  0  0  0]
 [ 0  0  0  0  0 37  0  0  0  0]
 [ 0  0  0  0  0  0 35  0  1  0]
 [ 0  0  0  0  0  0  0 36  0  0]
 [ 0  2  0  0  0  0  0  0 33  0]
 [ 0  0  0  0  0  0  0  1  0 35]]

```

Kernel: rbf

```

Accuracy : 0.9917
Precision: 0.9920
Recall    : 0.9915
F1-score  : 0.9916
Confusion Matrix:
[[36  0  0  0  0  0  0  0  0  0]
 [ 0 36  0  0  0  0  0  0  0  0]
 [ 0  0 35  0  0  0  0  0  0  0]
 [ 0  0  0 37  0  0  0  0  0  0]
 [ 0  0  0  0 36  0  0  0  0  0]
 [ 0  0  0  0  0 37  0  0  0  0]
 [ 0  0  0  0  0  0 36  0  0  0]
 [ 0  0  0  0  0  0  0 36  0  0]
 [ 0  2  0  0  0  0  0  0 33  0]
 [ 0  0  0  0  0  0  0  1  0 35]]

```

Kernel: sigmoid

```

Accuracy : 0.8861
Precision: 0.8925
Recall    : 0.8851
F1-score  : 0.8855
Confusion Matrix:
[[33  0  0  0  0  0  0  0  0  0]
 [ 0 33  0  0  0  0  0  0  0  0]
 [ 0  0 33  0  0  0  0  0  0  0]
 [ 0  0  0 33  0  0  0  0  0  0]
 [ 0  0  0  0 33  0  0  0  0  0]
 [ 0  0  0  0  0 33  0  0  0  0]
 [ 0  0  0  0  0  0 33  0  0  0]
 [ 0  0  0  0  0  0  0 33  0  0]
 [ 0  2  0  0  0  0  0  0 33  0]
 [ 0  0  0  0  0  0  0  1  0 35]]

```

```

from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split

```

```

data = load_digits()
X, y = data.data, data.target

```

```
#splits = [0.5, 0.4, 0.3, 0.2]
```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.4, random_state=42, stratify=y
)

```

```

from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix

```

```

kernels = ['linear', 'poly', 'rbf', 'sigmoid']
for k in kernels:

```

```

    model = SVC(kernel=k)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

```

```

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='macro') # precision
    rec = recall_score(y_test, y_pred, average='macro')    # recall
    f1 = f1_score(y_test, y_pred, average='macro')         # f1-score
    cm = confusion_matrix(y_test, y_pred)

```

```

    print(f"Kernel: {k}")
    print(f"Accuracy : {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall    : {rec:.4f}")
    print(f"F1-score  : {f1:.4f}")
    print("Confusion Matrix:\n", cm)
    print("-"*40)

```



```

Kernel: linear
Accuracy : 0.9694
Precision: 0.9711
Recall    : 0.9691
F1-score  : 0.9693
Confusion Matrix:
[[70  0  0  0  1  0  0  0  0  0]
 [ 0 72  0  0  0  0  0  0  0  1]
 [ 0  2 69  0  0  0  0  0  0  0]
 [ 0  0  0 71  0  0  0  0  0  2]

```

```
[ 0 0 0 0 72 0 0 0 0 0]
[ 0 0 0 0 0 72 0 0 0 1]
[ 0 1 0 0 0 0 71 0 0 0]
[ 0 0 0 0 0 0 0 72 0 0]
[ 0 7 0 0 0 2 1 0 60 0]
[ 0 0 0 0 0 0 0 3 1 68]]
```

Kernel: poly

Accuracy : 0.9861

Precision: 0.9866

Recall : 0.9860

F1-score : 0.9861

Confusion Matrix:

```
[[70 0 0 0 1 0 0 0 0 0]
 [ 0 73 0 0 0 0 0 0 0 0]
 [ 0 1 69 0 0 0 0 1 0 0]
 [ 0 0 0 73 0 0 0 0 0 0]
 [ 0 0 0 0 72 0 0 0 0 0]
 [ 0 0 0 0 0 72 0 0 0 1]
 [ 0 0 0 0 0 0 71 0 1 0]
 [ 0 0 0 0 0 0 0 72 0 0]
 [ 0 3 0 0 0 0 0 0 67 0]
 [ 0 0 0 0 0 0 0 2 0 70]]
```

Kernel: rbf

Accuracy : 0.9861

Precision: 0.9867

Recall : 0.9860

F1-score : 0.9861

Confusion Matrix:

```
[[70 0 0 0 1 0 0 0 0 0]
 [ 0 73 0 0 0 0 0 0 0 0]
 [ 0 1 70 0 0 0 0 0 0 0]
 [ 0 0 0 73 0 0 0 0 0 0]
 [ 0 0 0 0 72 0 0 0 0 0]
 [ 0 0 0 0 0 72 0 0 0 1]
 [ 0 0 0 0 0 0 72 0 0 0]
 [ 0 0 0 0 0 0 0 72 0 0]
 [ 0 3 0 0 0 0 0 0 67 0]
 [ 0 0 0 0 0 0 0 4 0 68]]
```

Kernel: sigmoid

Accuracy : 0.9026

Precision: 0.9075

Recall : 0.9021

F1-score : 0.9027

Confusion Matrix:

```
[[67 0 0 0 1 0 0 0 0 0]
 [ 0 67 0 0 0 0 0 0 0 0]
 [ 0 0 67 0 0 0 0 0 0 0]
 [ 0 0 0 67 0 0 0 0 0 0]
 [ 0 0 0 0 67 0 0 0 0 0]
 [ 0 0 0 0 0 67 0 0 0 0]
 [ 0 0 0 0 0 0 67 0 0 0]
 [ 0 0 0 0 0 0 0 67 0 0]
 [ 0 3 0 0 0 0 0 0 67 0]
 [ 0 0 0 0 0 0 0 4 0 68]]
```

```
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
```

```
data = load_digits()
X, y = data.data, data.target
```

```
#splits = [0.5, 0.4, 0.3, 0.2]
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.5, random_state=42, stratify=y
)
```

```
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
```

```
kernels = ['linear', 'poly', 'rbf', 'sigmoid']
for k in kernels:
```

```
    model = SVC(kernel=k)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
```

```
    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='macro') # precision
    rec = recall_score(y_test, y_pred, average='macro') # recall
    f1 = f1_score(y_test, y_pred, average='macro') # f1-score
    cm = confusion_matrix(y_test, y_pred)
```

```
    print(f"Kernel: {k}")
    print(f"Accuracy : {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall : {rec:.4f}")
    print(f"F1-score : {f1:.4f}")
    print("Confusion Matrix:\n", cm)
    print("-"*40)
```

```

Kernel: linear
Accuracy : 0.9733
Precision: 0.9736
Recall   : 0.9731
F1-score : 0.9731
Confusion Matrix:
[[88  0  0  0  1  0  0  0  0  0]
 [ 0 90  0  0  0  0  0  0  1  0]
 [ 0  1 87  0  0  0  0  0  0  0]
 [ 0  0  1 90  0  0  0  0  0  1]
 [ 0  0  0  0 91  0  0  0  0  0]
 [ 0  0  0  0  0 89  0  0  0  2]
 [ 0  0  0  0  0  0 90  0  1  0]
 [ 0  0  0  0  0  0  0 88  1  0]
 [ 0  6  0  1  0  0  1  0 79  0]
 [ 0  1  0  0  0  1  0  4  1 83]]

```

```

-----
Kernel: poly
Accuracy : 0.9822
Precision: 0.9825
Recall   : 0.9822
F1-score : 0.9822
Confusion Matrix:
[[88  0  0  0  1  0  0  0  0  0]
 [ 0 91  0  0  0  0  0  0  0  0]
 [ 0  1 86  0  0  0  0  1  0  0]
 [ 0  0  1 90  0  0  0  0  0  1]
 [ 0  0  0  0 91  0  0  0  0  0]
 [ 0  0  0  0  0 89  0  0  0  2]
 [ 0  0  0  0  0  0 90  0  1  0]
 [ 0  0  0  0  0  0  0 89  0  0]
 [ 0  2  0  0  0  0  0  0 85  0]
 [ 0  0  0  0  0  1  0  5  0 84]]

```

```

-----
Kernel: rbf
Accuracy : 0.9811
Precision: 0.9817
Recall   : 0.9810
F1-score : 0.9810
Confusion Matrix:
[[88  0  0  0  1  0  0  0  0  0]
 [ 0 91  0  0  0  0  0  0  0  0]
 [ 0  1 87  0  0  0  0  0  0  0]
 [ 0  0  0 91  0  0  0  0  1  0]
 [ 0  0  0  0 90  0  0  1  0  0]
 [ 0  0  0  0  0 90  0  0  0  1]
 [ 0  0  0  0  0  0 90  0  1  0]
 [ 0  0  0  0  0  0  0 89  0  0]
 [ 0  3  0  0  0  0  0  0 84  0]
 [ 0  0  0  0  0  2  0  6  0 82]]

```

```

-----
Kernel: sigmoid
Accuracy : 0.9066
Precision: 0.9106
Recall   : 0.9062
F1-score : 0.9072
Confusion Matrix:
[[86  0  0  0  3  0  0  0  0  0]

```

```

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

```

```

from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, label_binarize
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, confusion_matrix, roc_curve, auc)

```

```

# ----- LOAD DATA -----

```

```

digits = load_digits()
X, y = digits.data, digits.target
classes = np.unique(y)

```

```

# ----- PREPROCESS -----

```

```

scaler = StandardScaler()
X = scaler.fit_transform(X)

```

```

# ----- PARAMETERS -----

```

```

splits = [0.5, 0.6, 0.7, 0.8] # Train ratios
results = []

```

```

for split in splits:
    print(f"\n===== Train-Test Split: {int(split*100)}-{100-int(split*100)} =====\n")

```

```

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=1-split, stratify=y, random_state=42
)

# ----- MODEL -----
mlp = MLPClassifier(hidden_layer_sizes=(100,100),
                    solver="adam",
                    learning_rate_init=0.01,
                    max_iter=500,
                    random_state=42,
                    verbose=False) # training progress off for clarity

mlp.fit(X_train, y_train)

# ----- PREDICTION -----
y_pred = mlp.predict(X_test)

# ----- METRICS -----
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, average="macro")
rec = recall_score(y_test, y_pred, average="macro")
f1 = f1_score(y_test, y_pred, average="macro")

results.append([f"{int(split*100)}-{100-int(split*100)}", acc, prec, rec, f1])

print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall : {rec:.4f}")
print(f"F1-score : {f1:.4f}")

# ----- CONFUSION MATRIX -----
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(7,6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=classes, yticklabels=classes)
plt.title(f"Confusion Matrix (MLP) | Split={int(split*100)}-{100-int(split*100)}")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()

# ----- TRAINING LOSS CURVE -----
plt.figure(figsize=(6,5))
plt.plot(mlp.loss_curve_, label="Training Loss")
plt.title(f"Training Loss Curve (MLP) | Split={int(split*100)}-{100-int(split*100)}")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

# ----- ROC & AUC (Class 1 vs All only) -----
y_test_bin = label_binarize(y_test, classes=classes)
y_score = mlp.predict_proba(X_test)

plt.figure(figsize=(7,5))
fpr, tpr, _ = roc_curve(y_test_bin[:, 1], y_score[:, 1]) # class 1 vs all
roc_auc = auc(fpr, tpr)

plt.plot(fpr, tpr, label=f"Class 1 vs All (AUC={roc_auc:.2f})", color="blue")
plt.plot([0,1], [0,1], 'k--')

plt.title(f"ROC Curve (MLP) | Class 1 vs All | Split={int(split*100)}-{100-int(split*100)}")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.show()




# ----- COMPARISON TABLE -----
df = pd.DataFrame(results, columns=["Split", "Accuracy", "Precision", "Recall", "F1-score"])
print("\n==== Performance Comparison =====")
print(df)

```



===== Train-Test Split: 50-50 =====


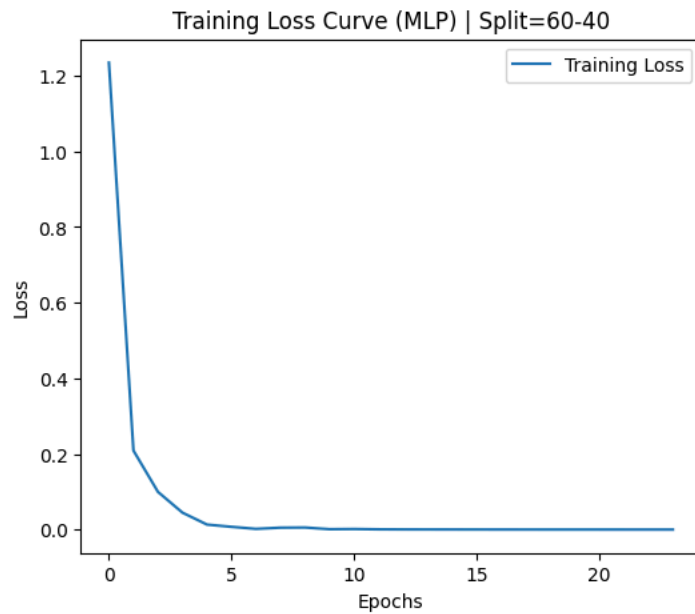

Accuracy : 0.9611
Precision: 0.9617
Recall : 0.9608
F1-score : 0.9609




0.0 0.2 0.4 0.6 0.8 1.0
False Positive Rate

===== Train-Test Split: 60-40 =====


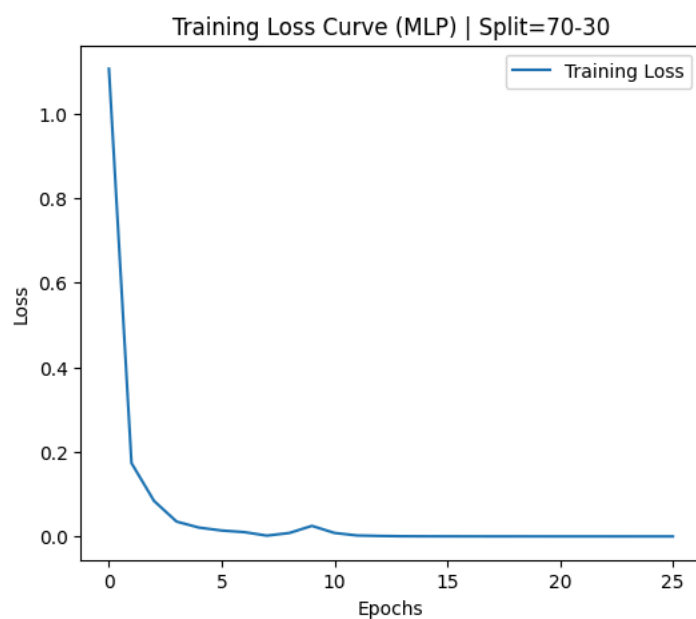
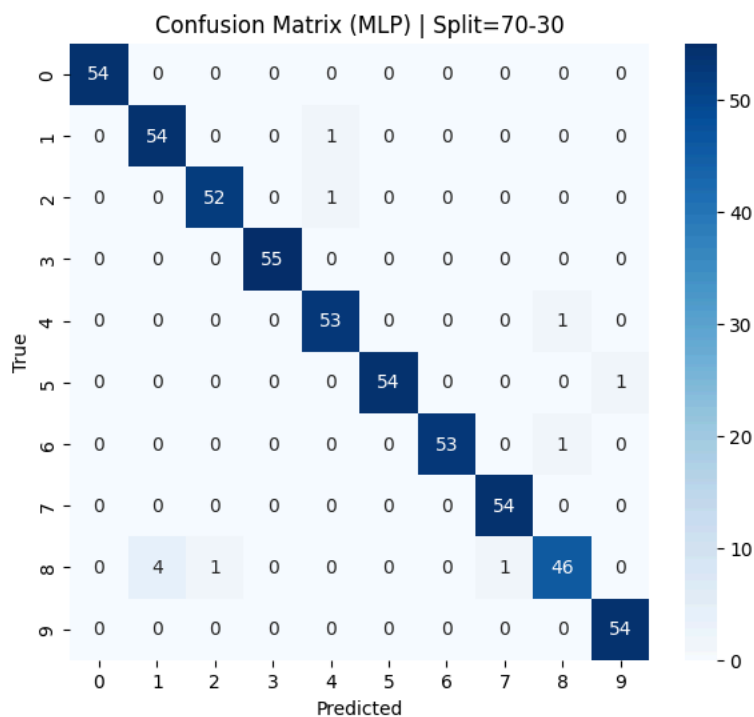
Accuracy : 0.9680
Precision: 0.9687
Recall : 0.9678
F1-score : 0.9679






===== Train-Test Split: 70-30 =====




Accuracy : 0.9796
Precision: 0.9798
Recall : 0.9792
F1-score : 0.9792






===== Train-Test Split: 80-20 =====

Accuracy : 0.9778
Precision: 0.9775
Recall : 0.9774
F1-score : 0.9771





==== Performance Comparison ====

	Split	Accuracy	Precision	Recall	F1-score
0	50-50	0.961068	0.961746	0.960812	0.960880
1	60-40	0.968011	0.968735	0.967795	0.967858
2	70-30	0.979630	0.979777	0.979235	0.979249
3	80-20	0.977778	0.977459	0.977381	0.977146

```

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, label_binarize
from sklearn.neural_network import MLPClassifier
from sklearn.multiclass import OneVsRestClassifier
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_curve, auc

# ----- LOAD DATA -----
digits = load_digits()
X, y = digits.data, digits.target
classes = np.unique(y)
n_classes = len(classes)

# ----- PREPROCESS -----
scaler = StandardScaler()
X = scaler.fit_transform(X)

# ----- PCA (95% variance) -----
pca = PCA(n_components=0.95, random_state=42)
X_pca = pca.fit_transform(X)
print(f"Original shape: {X.shape}, After PCA: {X_pca.shape}")

# ----- BEST SPLIT (70:30) -----
X_train, X_test, y_train, y_test = train_test_split(
    X_pca, y, test_size=0.3, stratify=y, random_state=42
)

# Binarize for ROC
y_train_bin = label_binarize(y_train, classes=classes)
y_test_bin = label_binarize(y_test, classes=classes)

# ----- MODEL -----
mlp = OneVsRestClassifier(
    MLPClassifier(hidden_layer_sizes=(100,100),
                  solver="adam",
                  learning_rate_init=0.01,
                  max_iter=5000,
                  random_state=42)
)

mlp.fit(X_train, y_train_bin)
y_pred_bin = mlp.predict(X_test)
y_score = mlp.predict_proba(X_test)

# Convert one-hot predictions back to class labels
y_pred = np.argmax(y_pred_bin, axis=1)

# ----- METRICS -----
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, average="macro")
rec = recall_score(y_test, y_pred, average="macro")
f1 = f1_score(y_test, y_pred, average="macro")

print("\n===== PCA + MLP (Digits, Best Split: 70-30) =====")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall : {rec:.4f}")
print(f"F1-score : {f1:.4f}")

# ----- CONFUSION MATRIX -----
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(7,6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Greens",
            xticklabels=classes, yticklabels=classes)
plt.title("Confusion Matrix (MLP + PCA, Digits, 70:30 Split)")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()

# ----- ROC & AUC (Only One-vs-All, e.g., Class 1 vs All) -----
target_class = 1 # change this to any digit (0-9) you want to test

fpr, tpr, _ = roc_curve(y_test_bin[:, target_class], y_score[:, target_class])
roc_auc = auc(fpr, tpr)

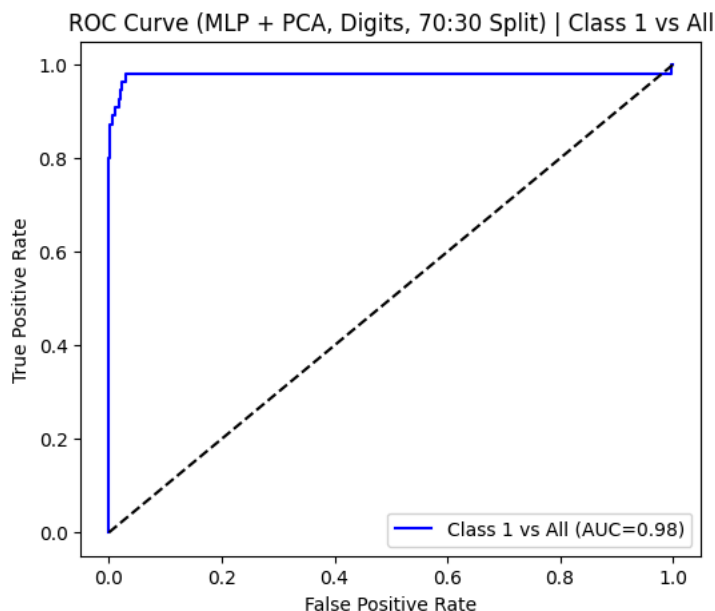

plt.figure(figsize=(6,5))

```

```
plt.plot(fpr, tpr, label=f"Class {target_class} vs All (AUC={roc_auc:.2f})", color="blue")
plt.plot([0,1], [0,1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title(f"ROC Curve (MLP + PCA, Digits, 70:30 Split) | Class {target_class} vs All")
plt.legend(loc="lower right")
plt.show()
```

Original shape: (1797, 64), After PCA: (1797, 40)

```
===== PCA + MLP (Digits, Best Split: 70-30) =====
Accuracy : 0.9500
Precision: 0.9579
Recall    : 0.9495
F1-score  : 0.9510
```



Start coding or [generate](#) with AI.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import label_binarize, StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.multiclass import OneVsRestClassifier
from sklearn.decomposition import PCA
```

```

import seaborn as sns

# ----- PARAMETERS -----
mom = 0.9
lr = 0.001
ep = 100
split = 0.5 # 50:50 split

# ----- DATA (Dummy for Example) -----
# Replace this with actual dataset (Wine/Digits)
X = np.random.rand(200, 10) # 200 samples, 10 features
y = np.random.randint(0, 3, 200) # 3-class labels (0,1,2)

# ----- PREPROCESS -----
scaler = StandardScaler()
X = scaler.fit_transform(X)

# PCA (reduce features, e.g. 95% variance retained)
pca = PCA(n_components=0.95)
X_pca = pca.fit_transform(X)
print(f"Original shape: {X.shape}, After PCA: {X_pca.shape}")

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_pca, y, test_size=split, random_state=42)

# Binarize labels (for multi-class handling)
classes = np.unique(y)
n_classes = len(classes)
y_train_bin = label_binarize(y_train, classes=classes)
y_test_bin = label_binarize(y_test, classes=classes)

# ----- MODEL -----
mlp_pca = OneVsRestClassifier(
    MLPClassifier(hidden_layer_sizes=(100,),
                  max_iter=ep,
                  learning_rate_init=lr,
                  momentum=mom,
                  solver='sgd',
                  random_state=42)
)

mlp_pca.fit(X_train, y_train_bin)
y_pred = mlp_pca.predict(X_test)

# ----- METRICS -----
# Convert predictions back to class labels
y_pred_classes = np.argmax(y_pred, axis=1)

acc = accuracy_score(y_test, y_pred_classes)
print(f"Accuracy with PCA: {acc:.2f}")

print("\nClassification Report (PCA):")
print(classification_report(y_test, y_pred_classes, target_names=[f"Class {c}" for c in classes]))

# ----- CONFUSION MATRIX -----
cm = confusion_matrix(y_test, y_pred_classes)

plt.figure(figsize=(5, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Greens", xticklabels=classes, yticklabels=classes)
plt.title("Confusion Matrix (MLP + PCA)")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()

```

Original shape: (200, 10), After PCA: (200, 10)
Accuracy with PCA: 0.29

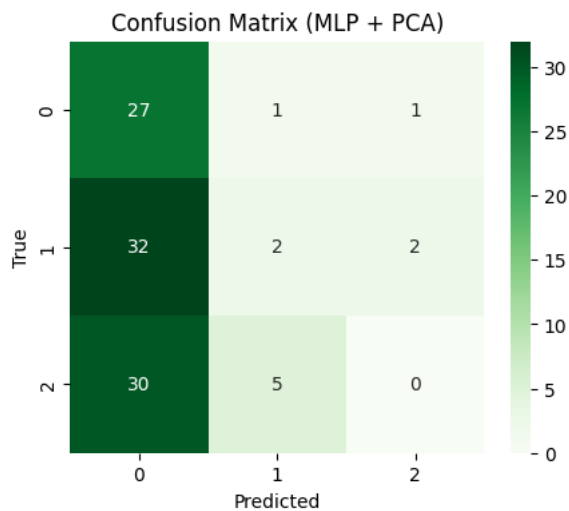
Classification Report (PCA):

	precision	recall	f1-score	support
Class 0	0.30	0.93	0.46	29
Class 1	0.25	0.06	0.09	36
Class 2	0.00	0.00	0.00	35
accuracy			0.29	100
macro avg	0.18	0.33	0.18	100
weighted avg	0.18	0.29	0.17	100

/usr/local/lib/python3.12/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimiz
warnings.warn(

/usr/local/lib/python3.12/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimiz
warnings.warn(

/usr/local/lib/python3.12/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimiz
warnings.warn(



Random Forest experiments: multiple splits, default + tuned, confusion heatmaps, ROC/AUC, optional PCA

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import (accuracy_score, precision_score, recall_score, f1_score,
                             confusion_matrix, roc_curve, auc, roc_auc_score)
from sklearn.preprocessing import StandardScaler, label_binarize
from sklearn.decomposition import PCA
from sklearn.multiclass import OneVsRestClassifier

# --- Configuration ---
splits = [0.5, 0.6, 0.7, 0.8]          # train ratios
do_grid_search = True                  # toggle tuned RF
rf_param_grid = {                      # simple grid for tuning (modify as needed)
    'n_estimators': [50, 100],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5]
}
random_state = 42
save_plots = False                    # set True to save figures instead of plt.show()

# number of classes
n_classes = len(np.unique(y))

# results collection
results = []

# optional: do PCA version as well
do_pca_versions = True
pca_variance = 0.95                  # keep components that explain 95% variance

# helper: plot + optionally save
def show_or_save(fig, title):
    if save_plots:
        fname = f"{title.replace(' ', '_')}.png"
        fig.savefig(fname, bbox_inches='tight')
    plt.close(fig)
```

```

        plt.show()
    else:
        plt.show()

# Main loop (first without PCA, then with PCA if requested)
for use_pca in [False, True] if do_pca_versions else [False]:
    version_tag = "PCA" if use_pca else "NoPCA"
    X_work = X.copy()
    # If using PCA, scale first (RF doesn't need scaling but PCA benefits)
    if use_pca:
        scaler = StandardScaler()
        X_scaled = scaler.fit_transform(X_work)
        pca = PCA(n_components=pca_variance, svd_solver='full', random_state=random_state)
        X_reduced = pca.fit_transform(X_scaled)
        print(f"[{version_tag}] PCA reduced dims: {X_reduced.shape[1]}")
        X_work = X_reduced
    else:
        # keep original X (no scaling needed for RF), but for ROC/AUC we won't scale
        pass

for split in splits:
    train_size = split
    test_size = 1 - split
    print(f"\n--- {version_tag} | Train-Test: {int(split*100)}:{int(test_size*100)} ---")

    X_train, X_test, y_train, y_test = train_test_split(
        X_work, y, train_size=train_size, random_state=random_state, stratify=y
    )

    # ---- 1) Default Random Forest ----
    rf_default = RandomForestClassifier(random_state=random_state)
    rf_default.fit(X_train, y_train)
    y_pred = rf_default.predict(X_test)
    y_proba = rf_default.predict_proba(X_test) # for ROC/AUC

    # metrics
    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='macro', zero_division=0)
    rec = recall_score(y_test, y_pred, average='macro', zero_division=0)
    f1 = f1_score(y_test, y_pred, average='macro', zero_division=0)
    cm = confusion_matrix(y_test, y_pred)

    # store
    results.append({
        'version': version_tag,
        'split': f"{int(split*100)}:{int(test_size*100)}",
        'model': 'RF_default',
        'acc': acc, 'precision': prec, 'recall': rec, 'f1': f1
    })

    # confusion heatmap
    fig = plt.figure(figsize=(5,4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
    plt.title(f"Confusion: RF_default | {version_tag} | Train-Test: {int(split*100)}:{int(test_size*100)}")
    plt.xlabel("Predicted"); plt.ylabel("True")
    show_or_save(fig, f"Confusion_RF_default_{version_tag}_Train-Test: {int(split*100)}:{int(test_size*100)}")

    # ROC & AUC (multiclass)
    # Binarize true labels
    y_test_bin = label_binarize(y_test, classes=np.arange(n_classes))
    # compute ROC per class
    fpr, tpr, roc_auc = dict(), dict(), dict()
    for i in range(n_classes):
        fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_proba[:, i])
        roc_auc[i] = auc(fpr[i], tpr[i])

    # macro AUC across classes
    try:
        overall_auc = roc_auc_score(y_test_bin, y_proba, average='macro', multi_class='ovr')
    except Exception:
        overall_auc = np.mean(list(roc_auc.values()))

    # plot multiclass ROC
    fig = plt.figure(figsize=(6,5))
    for i in range(n_classes):
        plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC={roc_auc[i]:.2f})")
    plt.plot([0,1],[0,1], 'k--')
    plt.title(f"ROC: RF_default | {version_tag} | Split {int(split*100)} | AUC_macro={overall_auc:.2f}")
    plt.xlabel("False Positive Rate"); plt.ylabel("True Positive Rate")
    plt.legend(loc='lower right', fontsize='small')
    show_or_save(fig, f"ROC_RF_default_{version_tag}_Train-Test: {int(split*100)}:{int(test_size*100)}")

    # ALSO add auc to results

```



```

results[-1].update({'auc_macro': overall_auc})

# ---- 2) Tuned Random Forest (GridSearchCV) ----
if do_grid_search:
    # Use a shallow GridSearch to save time; increase cv or grid if you want finer tuning
    gs = GridSearchCV(RandomForestClassifier(random_state=random_state),
                      rf_param_grid, cv=3, scoring='accuracy', n_jobs=-1, verbose=0)
    gs.fit(X_train, y_train)
    best = gs.best_estimator_
    print("RF GridSearch best params:", gs.best_params_)

    y_pred_t = best.predict(X_test)
    y_proba_t = best.predict_proba(X_test)

    acc_t = accuracy_score(y_test, y_pred_t)
    prec_t = precision_score(y_test, y_pred_t, average='macro', zero_division=0)
    rec_t = recall_score(y_test, y_pred_t, average='macro', zero_division=0)
    f1_t = f1_score(y_test, y_pred_t, average='macro', zero_division=0)
    cm_t = confusion_matrix(y_test, y_pred_t)

    results.append({
        'version': version_tag,
        'split': f"{int(split*100)}:{int(test_size*100)}",
        'model': 'RF_tuned',
        'acc': acc_t, 'precision': prec_t, 'recall': rec_t, 'f1': f1_t
    })

# confusion heatmap tuned
fig = plt.figure(figsize=(5,4))
sns.heatmap(cm_t, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title(f"Confusion: RF_tuned | {version_tag} | Train-Test: {int(split*100)}:{int(test_size*100)}")
plt.xlabel("Predicted"); plt.ylabel("True")
show_or_save(fig, f"Confusion_RF_tuned_{version_tag}_Train-Test: {int(split*100)}:{int(test_size*100)}")

# ROC tuned
fpr_t, tpr_t, roc_auc_t = dict(), dict(), dict()
for i in range(n_classes):
    fpr_t[i], tpr_t[i], _ = roc_curve(y_test_bin[:, i], y_proba_t[:, i])
    roc_auc_t[i] = auc(fpr_t[i], tpr_t[i])
try:
    overall_auc_t = roc_auc_score(y_test_bin, y_proba_t, average='macro', multi_class='ovr')
except Exception:
    overall_auc_t = np.mean(list(roc_auc_t.values()))

fig = plt.figure(figsize=(6,5))
for i in range(n_classes):
    plt.plot(fpr_t[i], tpr_t[i], label=f"Class {i} (AUC={roc_auc_t[i]:.2f})")
plt.plot([0,1],[0,1], 'k--')
plt.title(f"ROC: RF_tuned | {version_tag} | Train-Test: {int(split*100)}:{int(test_size*100)} | AUC_macro={overall_auc_t:.2f}")
plt.xlabel("False Positive Rate"); plt.ylabel("True Positive Rate")
plt.legend(loc='lower right', fontsize='small')
show_or_save(fig, f"ROC_RF_tuned_{version_tag}_Train-Test: {int(split*100)}:{int(test_size*100)}")

# add tuned auc
results[-1].update({'auc_macro': overall_auc_t})


# Convert results to DataFrame and show
results_df = pd.DataFrame(results)
display(results_df.sort_values(['version', 'split', 'model'], ascending=[True, True, True]))

```




--- NoPCA | Train-Test: 50:50 ---

Confusion: RF_default | NoPCA | Train-Test: 50:50




ROC: RF_default | NoPCA | Split 50 | AUC_macro=1.00





RF GridSearch best params: {'max_depth': None, 'min_samples_split': 2, 'n_estimators': 100}

Confusion: RF_tuned | NoPCA | Train-Test: 50:50




ROC: RF_tuned | NoPCA | Train-Test: 50:50 | AUC_macro=1.00






--- NoPCA | Train-Test: 60:40 ---

Confusion: RF_default | NoPCA | Train-Test: 60:40

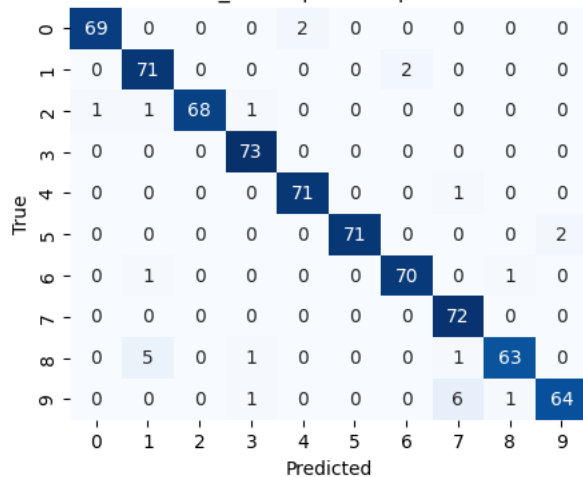


ROC: RF_default | NoPCA | Split 60 | AUC_macro=1.00




RF GridSearch best params: {'max_depth': 10, 'min_samples_split': 2, 'n_estimators': 50}

Confusion: RF_tuned | NoPCA | Train-Test: 60:40




ROC: RF_tuned | NoPCA | Train-Test: 60:40 | AUC_macro=1.00






--- NoPCA | Train-Test: 70:30 ---

Confusion: RF_default | NoPCA | Train-Test: 70:30





ROC: RF_default | NoPCA | Split 70 | AUC_macro=1.00




RF GridSearch best params: {'max_depth': None, 'min_samples_split': 2, 'n_estimators': 100}

Confusion: RF_tuned | NoPCA | Train-Test: 70:30






ROC: RF_tuned | NoPCA | Train-Test: 70:30 | AUC_macro=1.00




--- NoPCA | Train-Test: 80:19 ---

Confusion: RF_default | NoPCA | Train-Test: 80:19




ROC: RF_default | NoPCA | Split 80 | AUC_macro=1.00




RF GridSearch best params: {'max_depth': 10, 'min_samples_split': 2, 'n_estimators': 100}

Confusion: RF_tuned | NoPCA | Train-Test: 80:19




ROC: RF_tuned | NoPCA | Train-Test: 80:19 | AUC_macro=1.00




[PCA] PCA reduced dims: 40


--- PCA | Train-Test: 50:50 ---

Confusion: RF_default | PCA | Train-Test: 50:50




ROC: RF_default | PCA | Split 50 | AUC_macro=1.00






RF GridSearch best params: {'max_depth': 10, 'min_samples_split': 2, 'n_estimators': 100}

Confusion: RF_tuned | PCA | Train-Test: 50:50




ROC: RF_tuned | PCA | Train-Test: 50:50 | AUC_macro=1.00





--- PCA | Train-Test: 60:40 ---

Confusion: RF_default | PCA | Train-Test: 60:40





ROC: RF_default | PCA | Split 60 | AUC_macro=1.00






RF GridSearch best params: {'max_depth': None, 'min_samples_split': 5, 'n_estimators': 50}




--- PCA | Train-Test: 70:30 ---



8	0	4	0	2	0	2	0	1	43	0
9	0	0	0	0	0	1	1	1	0	51
	0	1	2	3	4	5	6	7	8	9

Predicted




RF GridSearch best params: {'max_depth': 10, 'min_samples_split': 5, 'n_estimators': 100}

Confusion: RF_tuned | PCA | Train-Test: 70:30

0	52	0	0	0	1	0	1	0	0	0
1	0	54	1	0	0	0	0	0	0	0
2	0	0	52	0	0	0	0	1	0	0
3	0	0	0	51	0	3	0	0	1	0
4	0	0	0	0	53	0	0	1	0	0
5	0	0	0	0	0	52	0	0	0	3
6	0	1	0	0	0	0	52	0	1	0
7	0	0	0	0	0	0	0	54	0	0
8	0	2	1	1	0	1	0	0	47	0
9	0	0	0	0	0	1	1	1	3	48
	0	1	2	3	4	5	6	7	8	9

True

Predicted



--- PCA | Train-Test: 80:19 ---

Confusion: RF_default | PCA | Train-Test: 80:19

0	35	0	0	0	0	0	1	0	0	0
---	----	---	---	---	---	---	---	---	---	---

1	0	35	0	0	0	1	0	0	0	0
2	0	0	35	0	0	0	0	0	0	0
3	0	0	1	36	0	0	0	0	0	0
4	0	0	0	0	34	0	0	2	0	0
5	0	0	0	0	0	37	0	0	0	0
6	0	1	0	0	0	0	34	0	1	0
7	0	0	0	0	0	0	0	36	0	0
8	0	3	1	0	0	0	0	1	29	1
9	0	0	0	0	0	1	1	1	0	33
	0	1	2	3	4	5	6	7	8	9

